

Training loss decreases; validation loss stays flat
(typical of low signal-to-noise regimes)

